

Binary Heaps & Priority Queues

- Priority Queue ADT
- Binary Heaps

How do you read your e-mails?

Oldest first?

Newest first?

Most important first?

What kinds of data structures model these methods?



What kinds of data structures model these methods?

Oldest first: FIFO (Queue)

Newest first: LIFO (Stack)

Most important first: **Priority Queue**

Priority Queue ADT

- Enforces some kind of “priority” on the items
- For general purposes, call the highest priority the “maximum” (or the minimum)
- Operations:
 - insert (**insert**)
 - remove maximum (**delMax**) xor remove minimum (**delMin**)
 - Note: You would not have both delMax and delMin in a single PQ—you would have one or the other depending on whether it is a Max PQ or a Min PQ
- Often, we store (key, value) pairs where the key is the “priority”
- **Key**: used to “rank” items—must be comparable

Recall Project 1.

Basically, you got the “top M ” integers by using a min PQ.

Priority Queue API

```
public class MaxPQ<Key extends Comparable<Key>>
```

```
    MaxPQ() create a priority queue
```

```
    MaxPQ(int max) create a priority queue of initial capacity max
```

```
    MaxPQ(Key[] a) create a priority queue from the keys in a[]
```

```
    void insert(Key v) insert a key into the priority queue
```

```
    Key max() return the largest key
```

```
    Key delMax() return and remove the largest key
```

```
    boolean isEmpty() is the priority queue empty?
```

```
    int size() number of keys in the priority queue
```

API for a generic priority queue

NOTE: a **MinPQ** would be basically the same except you would have **min()** to get the smallest key and **delMin()** to return and remove the smallest key.

How would you implement
a Priority Queue?

Unordered Array

insert : $O(1)$
delMin : $O(n)$

insert 23, 10, 17, 28, 34, 89, 22, 10
remove min

23	10	17	28	34	89	22	10
----	----	----	----	----	----	----	----

Ordered Array

insert: $O(N)$
delMin: $O(1)$

insert 23, 10, 17, 28, 34, 89, 22, 10
remove min

89	34	28	23	22	17	10	10
----	----	----	----	----	----	----	---------------

Key Operations: **insert** & **delMax** (or **delMin**)

Option 1

Unordered array

Worst-case runtime analysis:

- **insert**
- **delMax**

Option 2

Ordered array

Worst-case runtime analysis:

- **insert**
- **delMax**

Key Operations: **insert** & **delMax** (or **delMin**)

Option 1

Unordered array

Worst-case runtime analysis:

- **insert: $O(1)$**
- **delMax: $O(n)$**

Option 2

Ordered array

Worst-case runtime analysis:

- **insert: $O(n)$**
- **delMax: $O(1)$**

Key Operations: insert & delMax (or delMin)

Option 1

Unordered array

Worst-case runtime analysis:

- **insert: $O(1)$**
- **delMax: $O(n)$**

LAZY!

Option 2


Ordered array

Worst-case runtime analysis:

- **insert: $O(n)$**
- **delMax: $O(1)$**

EAGER!

SPOILER ALERT!!! There's a better way...




data structure	insert	remove maximum
<i>ordered array</i>	N	1
<i>unordered array</i>	1	N
<i>heap</i>	$\log N$	$\log N$
<i>impossible</i>	1	1

Order of growth of worst-case running time for priority-queue implementations

Imagine a company...

- New people are always hired at the lowest level
- Only the CEO ever quits
- Final ranking is determined by an IQ test where any given employee can only be managed by another employee with a higher IQ.
- Any given person can only manage two people.


(insert) swim = bubble up
(delMax) Sink → bubble down



So...what's a (max binary) heap?

Definition. A binary tree is *heap-ordered* if the key in each node is larger than or equal to the keys in that node's two children (if any).


NOTE: This is max-ordered binary heap.
How would it be different if it was a min-ordered binary heap?



A heap-ordered complete binary tree

So...what's a (binary) heap?



Definition. A binary tree is *heap-ordered* if the key in each node is larger than or equal to the keys in that node's two children (if any).



A heap-ordered complete binary tree

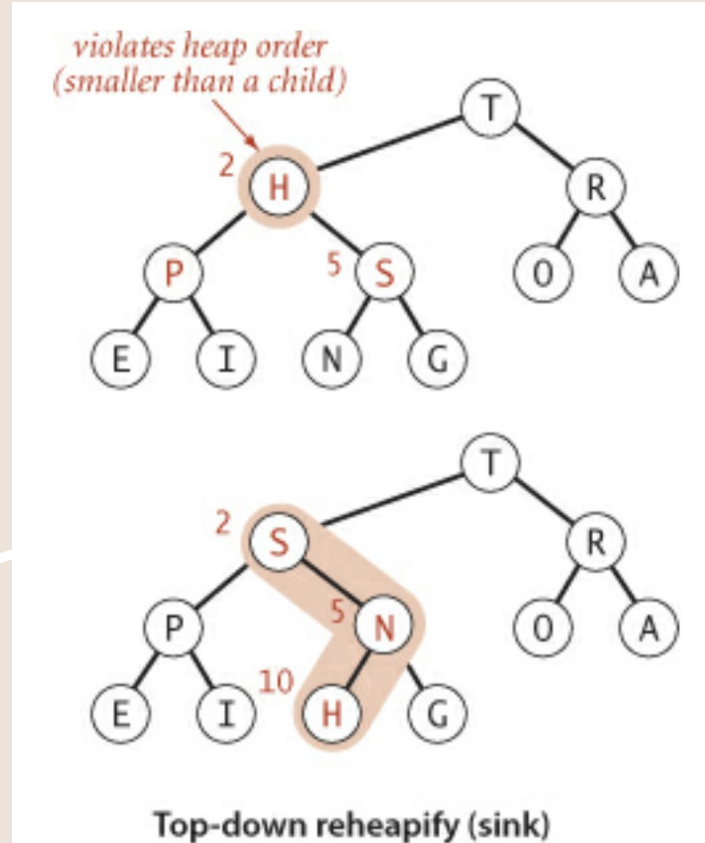
Where is the maximum? *ROOT*
Where would we insert a new node?
So...which locations should be easy to access?

Swimming up...






Bottom-up reheapify (swim)




Sinking down...



insert



remove the maximum



Heap operations


So why is it $O(\log N)$?

height : $O(\log N)$

complete tree
 \Rightarrow balanced tree

data structure	insert	remove maximum
<i>ordered array</i>	N	1
<i>unordered array</i>	1	N
<i>heap</i>	$\log N$	$\log N$
<i>impossible</i>	1	1

**Order of growth of worst-case running time
for priority-queue implementations**



A heap-ordered complete binary tree


So why is it $O(\log N)$?

data structure	insert	remove maximum
<i>ordered array</i>	N	1
<i>unordered array</i>	1	N
<i>heap</i>	$\log N$	$\log N$
<i>impossible</i>	1	1



Order of growth of worst-case running time for priority-queue implementations

- A binary heap is always a complete binary tree,
- which means it is always balanced,
- which means the height is always $O(\log N)$,
- which means that in the worst case, you have to *swim* or *sink* an item on a path of length $O(\log N)$.

So how would you implement a binary heap?



A heap-ordered complete binary tree




Heap representations

- If a node is at index k , its children are at $2k$ and $2k + 1$
- If a node is at index k , its parent is at $\text{floor}(k/2)$
- Keep track of an insert index.
- Root (maximum) should always be at index 1
- Swim/Sink as appropriate


In the worst case, any path from root to node is $O(\log n)$.

root



Note that in this picture, index 0 is not used. But that is not a requirement. You could use index 0 as the root. It would just change the relationship between parents and children.

A node at index k would have children at $2k+1$ and $2k+2$.



Heap representations

Array-based Binary Heap $\left\lfloor \frac{n-1}{2} \right\rfloor$

parent: k

children: $2k+1, 2k+2$

← c

$k=0 \rightarrow 1, 2$

insert 23, 10, 17, 28, 34, 89, 22, 10
remove min

0

1

2

3

4

5

6

7

10

23

17

28

34

89

22

Common Misconceptions

- A binary heap is not the same thing as a priority queue. A priority queue is an ADT that stores and processes data according to some kind of priority ranking. A binary heap (which can be max-ordered or min-ordered) is a common way of implementing a PQ because it is a structure that does *insert* and *delMax* (or *delMin*) efficiently.
- A binary heap, in turn, can be implemented either as a tree or as an array.
- You can store anything in a PQ as long as you can assign some kind of ranking to it. Often, the “priority” is treated as the key in a key-value pair. e.g. (IQ, Employee), but the “key” and the “value” could also be the same.

What are some applications that could use a PO?

References

- [1] *Algorithms, Fourth Edition*; Robert Sedgewick and Kevin Wayne (and associated slides)
- [2] Book slides from Goodrich and Tamassia