

# Binary Search Trees

*An ordered symbol table*

A large, dark blue, curved shape that starts from the bottom left and extends diagonally upwards towards the right, filling the lower half of the slide.

# Intuition for Binary Search Tree


Combine the useful elements of linked lists with the useful elements of ordered arrays:

- Linked lists don't require resizing, making it easier (more flexibility) to add items
- Ordered arrays allow for binary search, making it easier (higher efficiency) to find items (also supports ordered operations)

**Solution: Binary Search Tree!**

# Binary Search Tree


- Recursive definition: A Binary Search Tree is a **null** link or a **Node** with a **left link** and a **right link** that each point to Binary Search Trees
- A Node has one parent (except the root) and two children (left and right)



Anatomy of a binary tree



# Binary Search Tree

- A Node can have a left child and a right child.
- Nodes are ordered such that any Node is greater than any element in its left subtree and less than any element in its right subtree.
- So an inorder traversal of a BST would yield the keys in sorted order.



**Anatomy of a binary search tree**

node count N




Two BSTs that represent  
the same set of keys


Same keys, different trees...

HOW DOES THIS HAPPEN?

node count N



Resulting Binary Trees depend on the order in which elements are inserted, which means the same keys can be organized in very different trees.




Two BSTs that represent  
the same set of keys

**Example:** *insert 0, insert 1, insert 2, insert 3, ..., insert 9*






**Example:** insert 4, insert 2, insert 6, insert 1, insert 3,  
insert 5, insert 7, insert 0, insert 8, insert 9







### successful search for R






### unsuccessful search for T



link is null  
so T is not in tree  
(search miss)



Search hit (left) and search miss (right) in a BST

inserting L




- Search for key first
- If it's found, reset the value
- If it isn't found, add a new Node at the null link

Insertion into a BST




BST trace for standard indexing client


best case



typical case



worst case



BST possibilities

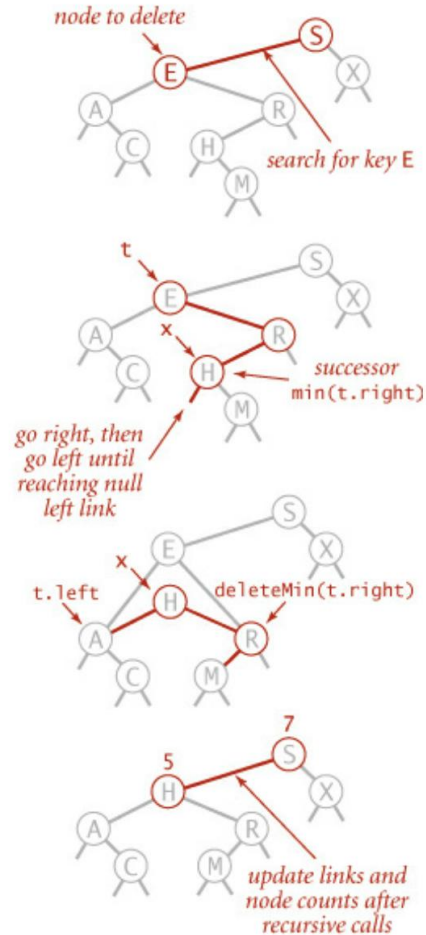
algorithm (data structure)	worst-case cost (after $N$ inserts)		average-case cost (after $N$ random inserts)		efficiently support ordered operations?
	search	insert	search hit	insert	
<i>sequential search</i> (unordered linked list)	$N$	$N$	$N/2$	$N$	no
<i>binary search</i> (ordered array)	$\lg N$	$N$	$\lg N$	$N/2$	yes
<i>binary tree search</i> (BST)	$N$	$N$	$1.39 \lg N$	$1.39 \lg N$	yes

Cost summary for basic symbol-table implementations (updated)

The good news is that the typical case is closer to the best case than the worst case!

# BST Deletion

### deleting E



Deletion in a BST

- Option 1: key doesn't exist in tree  
do nothing
- Option 2: key exists and the node has no children  
delete it
- Option 3: key exists and the node has one child  
make parent pointer point to node's child
- Option 4: key exists and the node has two children  
replace node w/ successor  
delete successor


**Example:** insert 4, insert 2, insert 6, insert 1, insert 3,  
insert 5, insert 7, insert 0, insert 8, insert 9, ~~delete 9, delete 7,~~  
~~delete 4, delete 2, delete 3, delete 6, delete 0, delete 1, delete 5,~~  
delete 8

# Analysis Summary

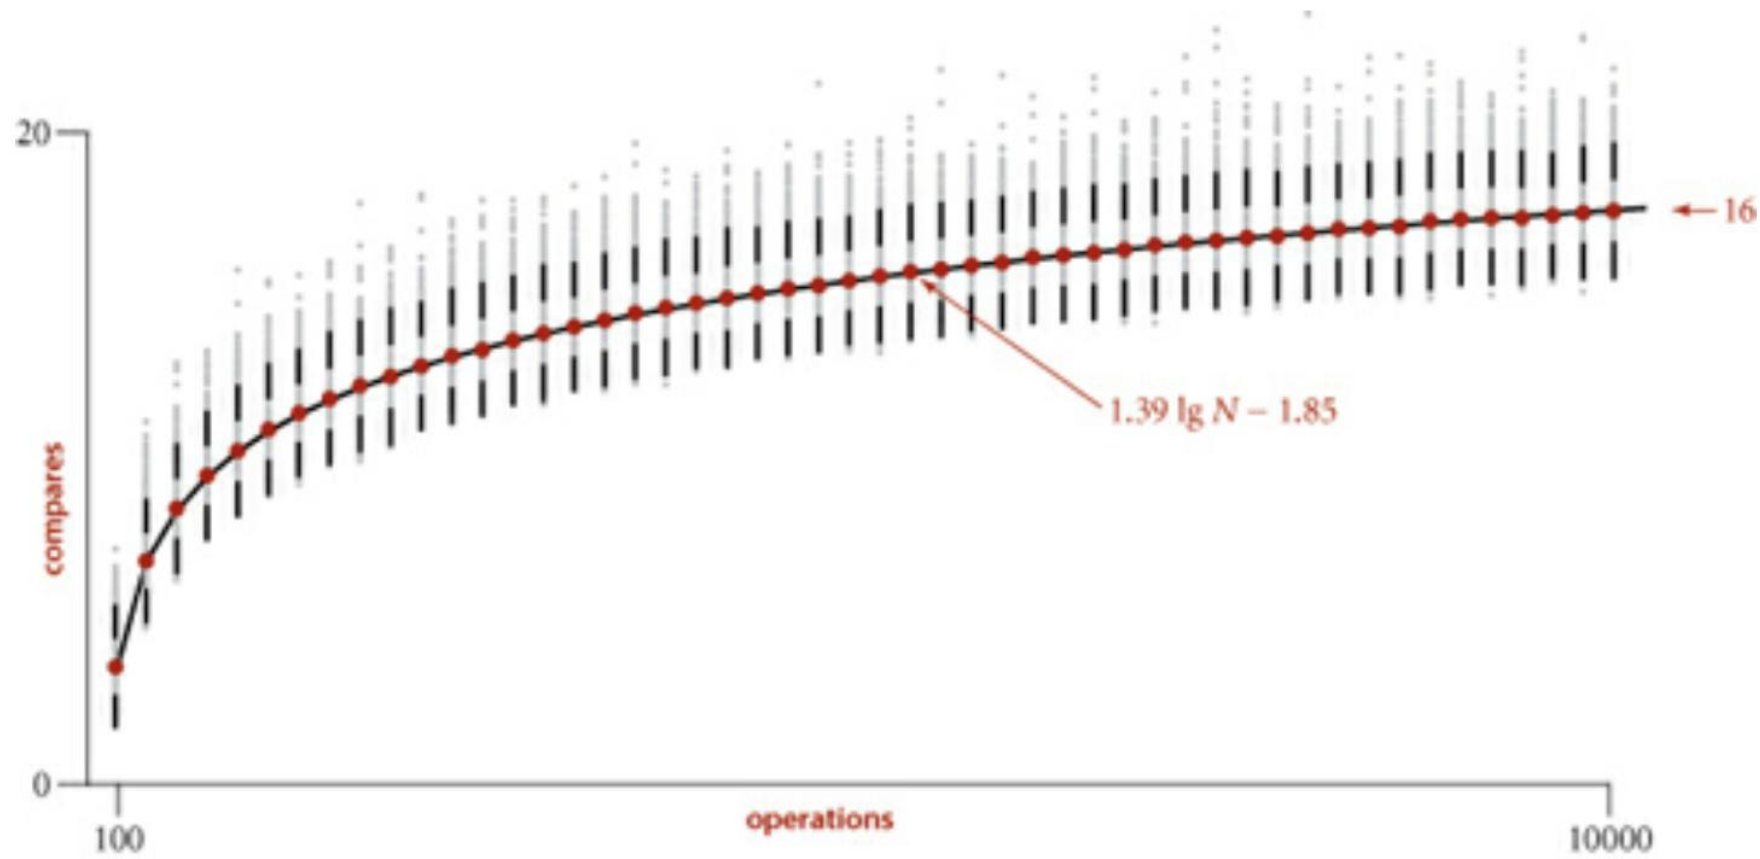
algorithm (data structure)	worst-case cost (after $N$ inserts)		average-case cost (after $N$ random inserts)		efficiently support ordered operations?
	search	insert	search hit	insert	
<i>sequential search (unordered linked list)</i>	$N$	$N$	$N/2$	$N$	no
<i>binary search (ordered array)</i>	$\lg N$	$N$	$\lg N$	$N/2$	yes
<i>binary tree search (BST)</i>	$N$	$N$	$1.39 \lg N$	$1.39 \lg N$	yes

Cost summary for basic symbol-table implementations (updated)





Typical BST, built from 256 random keys



Average path length to a random node in a BST built from random keys

# Binary Search Trees

## Pros

- Easy to implement
- Good average performance for basic operations

## Cons

- Bad worst case performance for basic operations
- WHY?

How can we improve upon this?

How can we improve upon this?

ENFORCE BALANCE

# References

- [1] *Algorithms, Fourth Edition*; Robert Sedgewick and Kevin Wayne (and associated slides)
- [2] Slides from <https://www.cs.princeton.edu/~rs/talks/LLRB/RedBlack.pdf>