

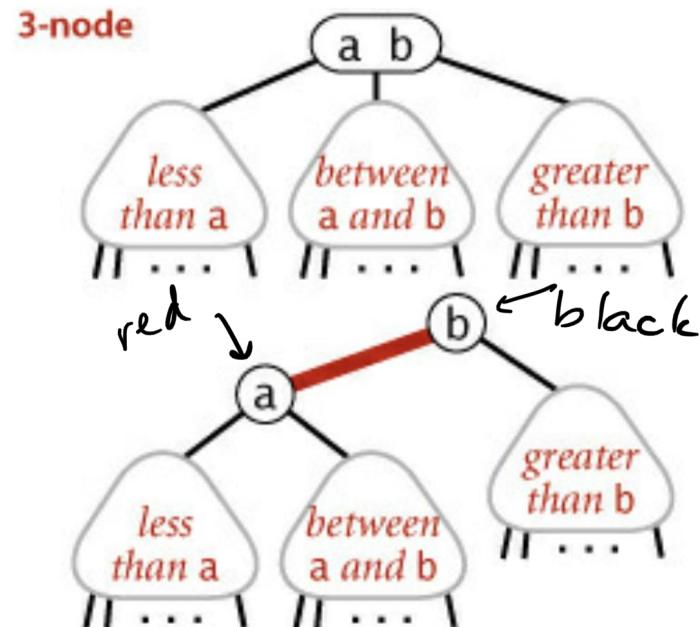
# Left-Leaning Red-Black Trees

# Left-Leaning Red-Black Trees

- An *ordered* symbol table
- A self-balancing tree
- Combines the simple implementation of a BST with the balancing ideas of a 2-3 tree
- Provides a tree implementation with a 1-to-1 correspondence with the more complex 2-3 tree

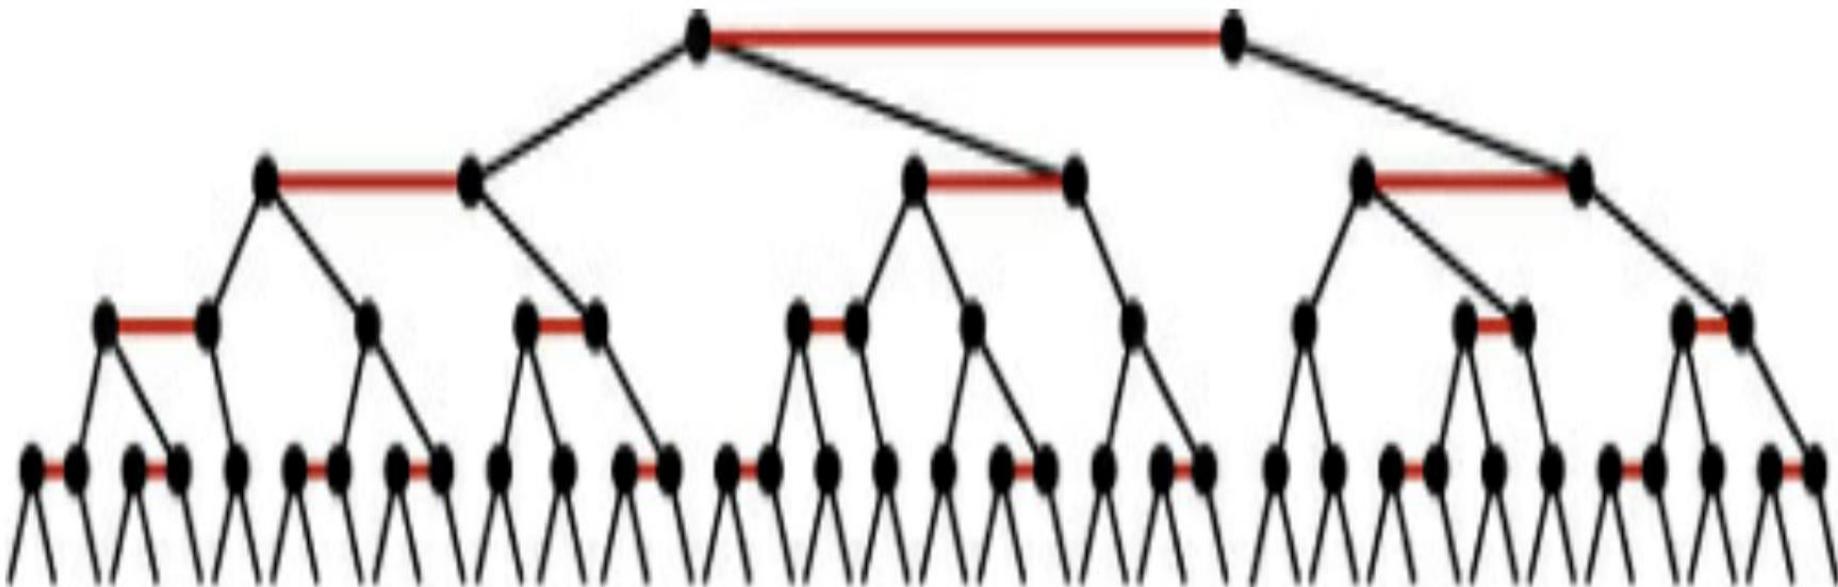
# (Left-leaning) Red-Black Trees

- Representation for 2-3 trees
- Basic Idea: Start with standard BST (2-nodes only) and add extra information to the nodes to represent 3-nodes
- 2 kinds of links:
  - red—bind together two nodes to represent 3-nodes
  - black—bind together the 2-3 tree
- 3-nodes are encoded as two nodes connected by a single left-leaning red link
- NOTE: We don't really encode "links." We just indicate a node's link to its parent by coloring the node itself.



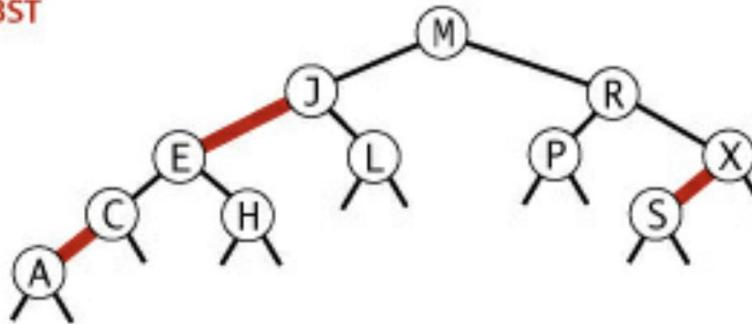
Encoding a 3-node with two 2-nodes  
connected by a left-leaning red link

“Flattening” the red links into horizontal links essentially turns the binary R-B tree into a 2-3 tree.



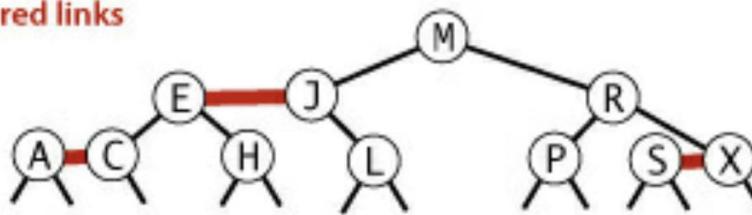
A red-black tree with horizontal red links is a 2-3 tree

red-black BST

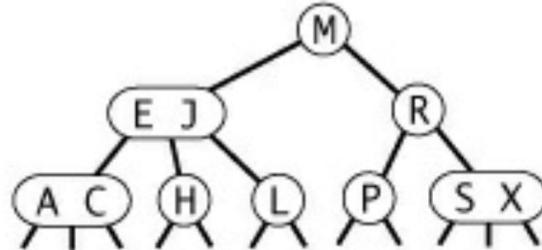


“Flattening” the red links into horizontal links essentially turns the binary R-B tree into a 2-3 tree.

horizontal red links



2-3 tree



1-1 correspondence between red-black BSTs and 2-3 trees

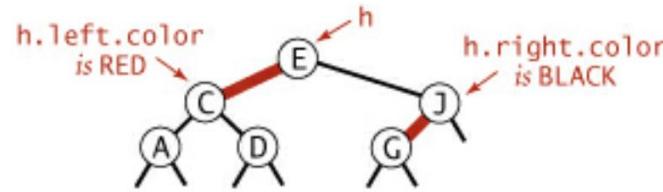
# Red-Black Tree: A BST with these properties...

- Links are **red** or **black**
  - Typically, we call the nodes red or black depending on the color of the link that connects it to its parent because there is no “link” object. It’s just a pointer from one node to another.
  - We can encode this by adding a boolean variable to each node (**true if red, false if black**).
- ~~\*A black node has *at most* one red child. (i.e. red nodes cannot have red siblings)~~
- ~~\*Red nodes have black children. (i.e. red nodes cannot have red children)~~
- ~~The root is always black.~~
- Tree has perfect *black balance* (all paths from the root to a leaf have the same number of black links—also known as the **black height** of the tree)

In a *left-leaning* red-black tree, all red links lean to the left. This is an implementation decision that can reduce the number of “cases” and therefore shorten the code.

# Color Representation

Each node has a new boolean variable for color (according to the link connecting it to its parent)—**true if red, false if black.**



```
private static final boolean RED = true;
private static final boolean BLACK = false;

private class Node
{
    Key key;           // key
    Value val;         // associated data
    Node left, right; // subtrees
    int N;             // # nodes in this subtree
    boolean color;    // color of link from
                      // parent to this node

    Node(Key key, Value val, int N, boolean color)
    {
        this.key = key;
        this.val = val;
        this.N = N;
        this.color = color;
    }
}

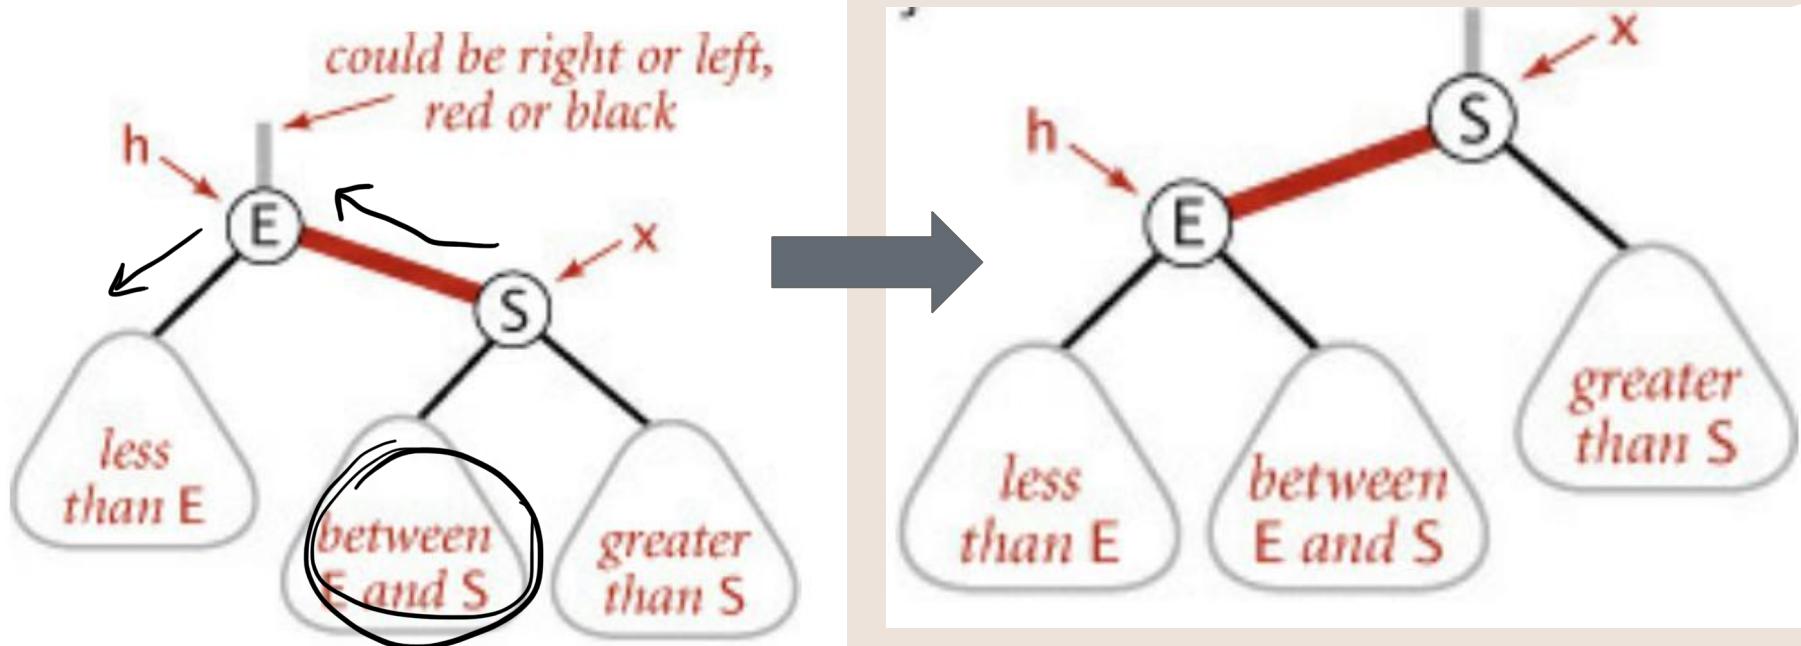
private boolean isRed(Node x)
{
    if (x == null) return false;
    return x.color == RED;
}
```

# Transformations

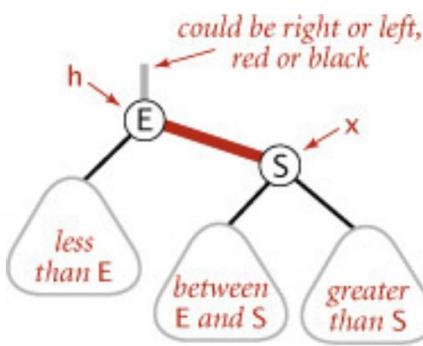
This is how we maintain the properties that ensure black height balance.



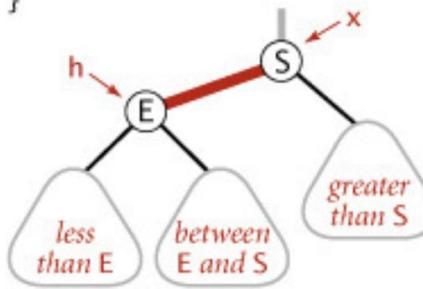
# Left Rotation



# Left Rotation Code

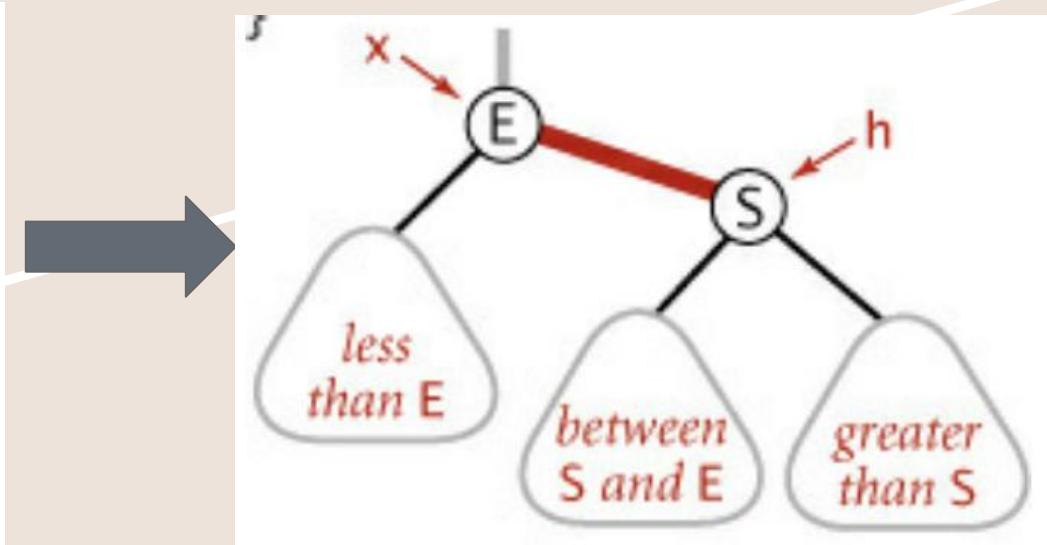
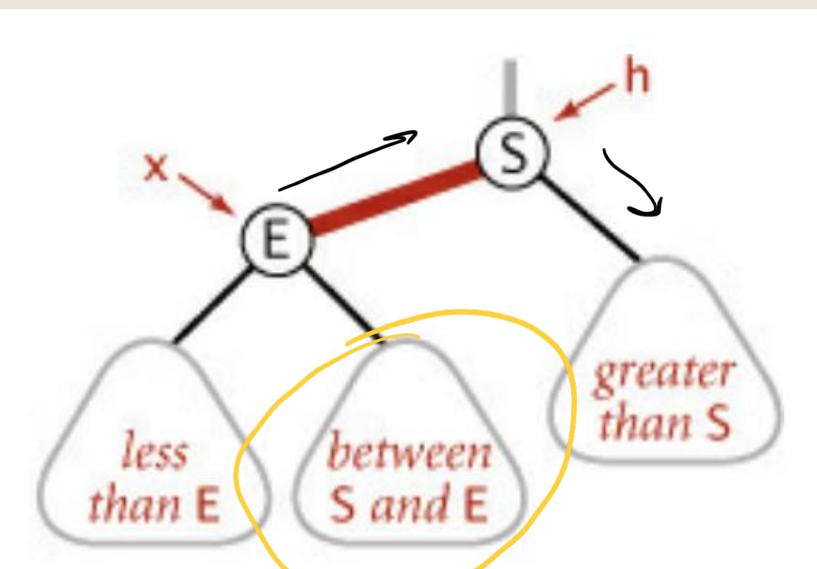


```
Node rotateLeft(Node h)
{
    Node x = h.right;
    h.right = x.left;
    x.left = h;
    x.color = h.color;
    h.color = RED;
    x.N = h.N;
    h.N = 1 + size(h.left)
        + size(h.right);
    return x;
}
```

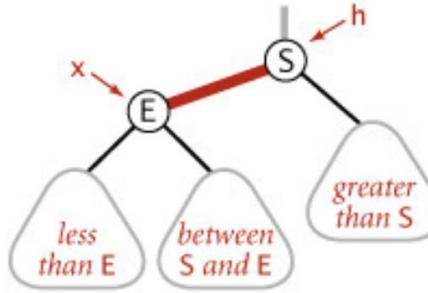


Left rotate (right link of h)

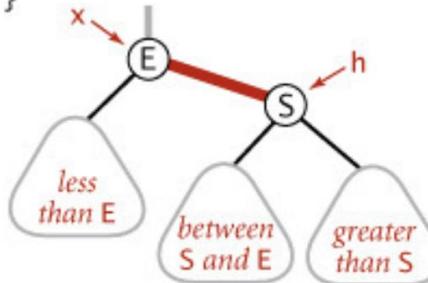
# Right Rotation



# Right Rotation Code

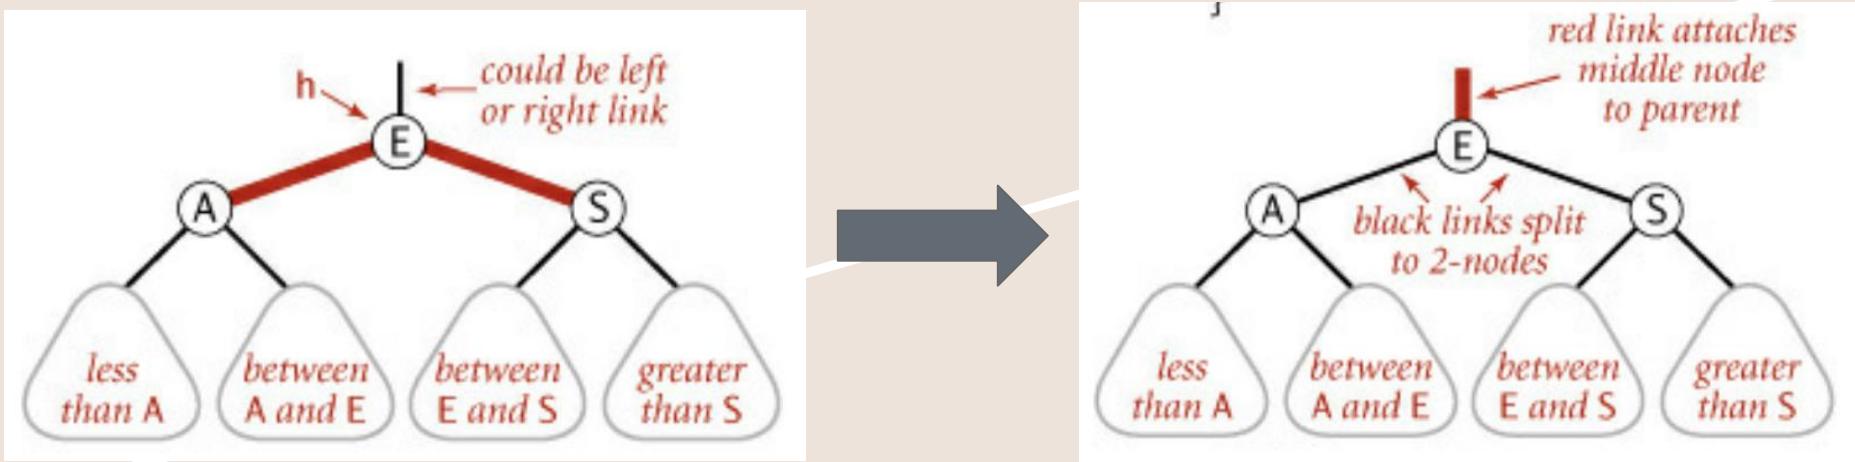


```
Node rotateRight(Node h)
{
    Node x = h.left;
    h.left = x.right;
    x.right = h;
    x.color = h.color;
    h.color = RED;
    x.N = h.N;
    h.N = 1 + size(h.left)
          + size(h.right);
    return x;
}
```

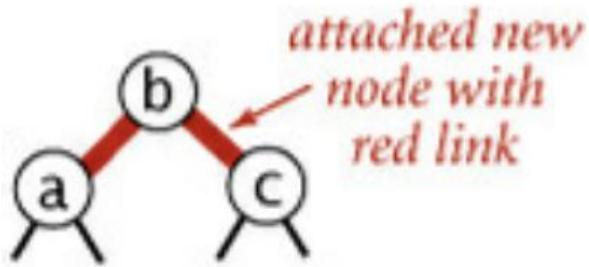


Right rotate (left link of h)

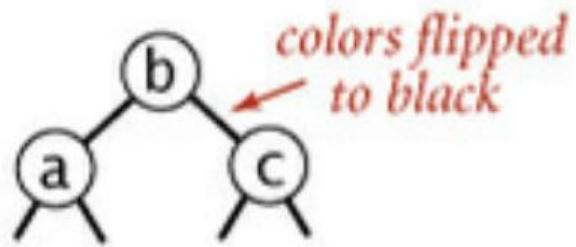
# Color Flip (equivalent to splitting a 4-node)



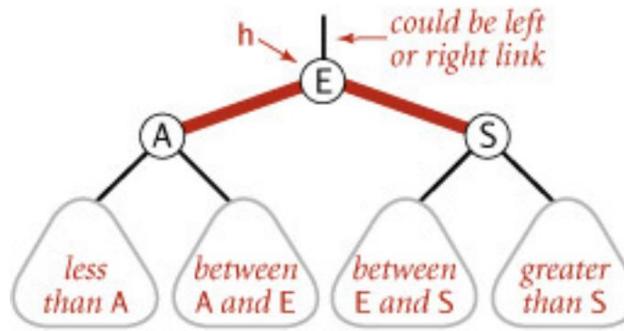
# Color Flip at the root



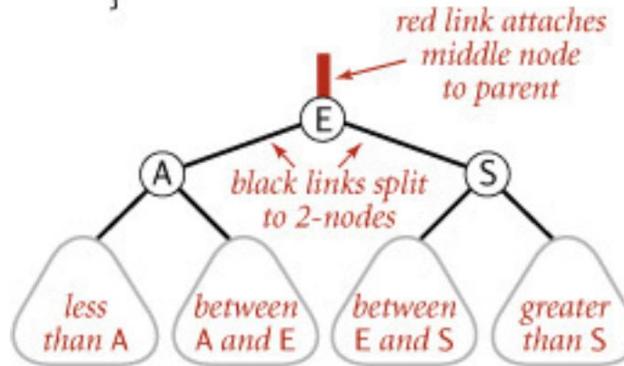
This is when the black height increases (like splitting a 4-node at the root).



# Color Flip Code



```
void flipColors(Node h)
{
    h.color = RED;
    h.left.color = BLACK;
    h.right.color = BLACK;
}
```



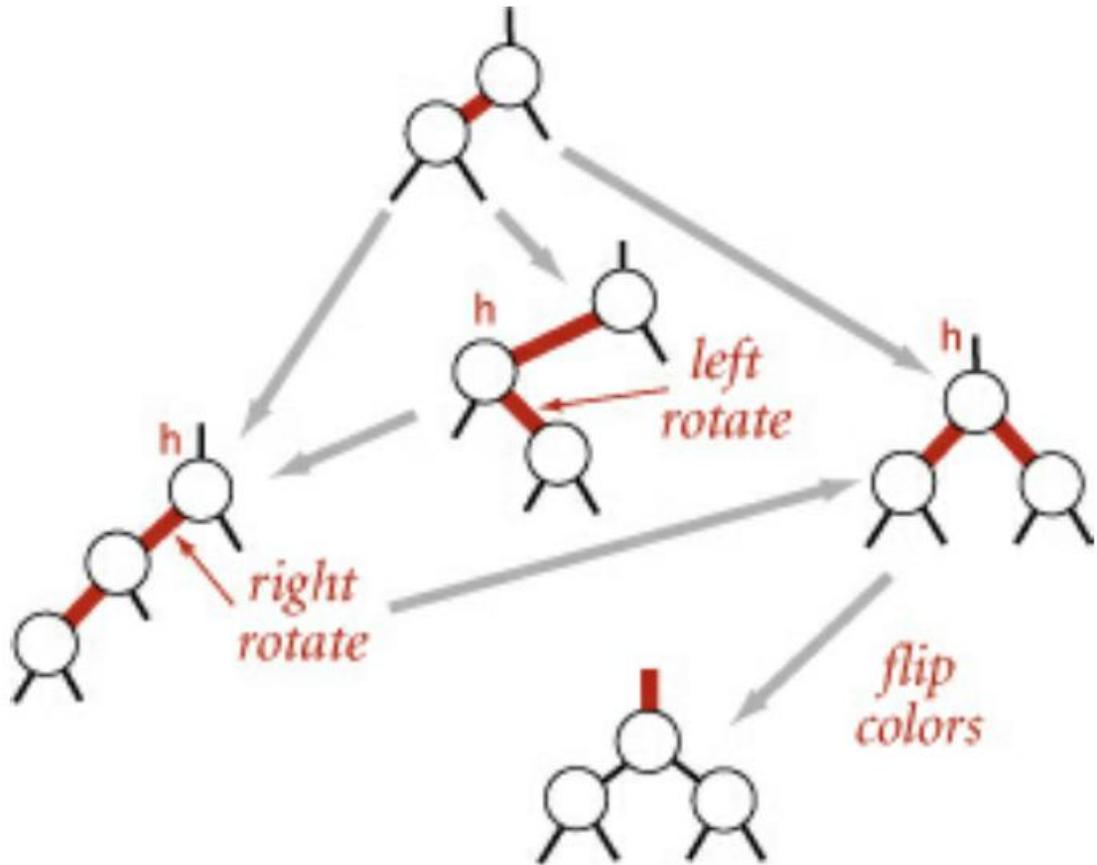
Flipping colors to split a 4-node

# Operations

put

- New nodes are always inserted at the bottom of the tree.
- New nodes are always red.
- Transformations are used to re-balance the black height of the tree by ensuring the properties discussed previously.
- Transformations may pass red links up the tree just like splitting 4-nodes.

# Fixing the tree...



Passing a red link up a red-black BST

- 3 main operations to maintain balance:
- rotate left
  - rotate right
  - color flip

Insertion uses those three operations as you “pass the red link up the tree”

Case 1:  $X.left$  is red,  $X.right$  is null or black.



Case 2:  $X.right$  is red,  $X.left$  is null or black.



Case 3:  $X.left$  is red,  $X.right$  is red. (Note that  $X$  itself will always be black in this case.)



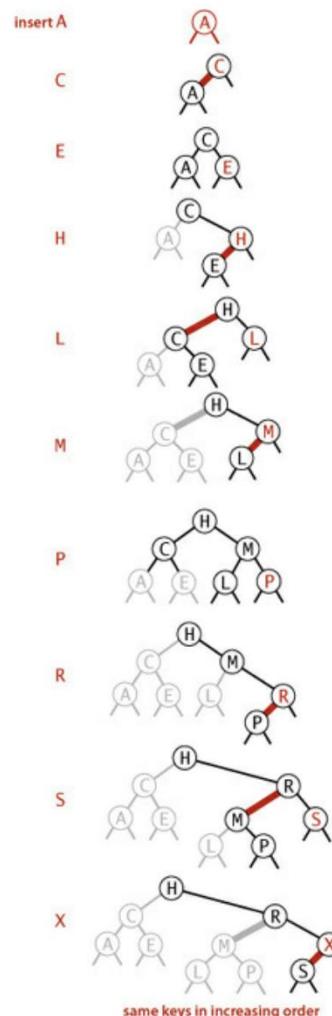
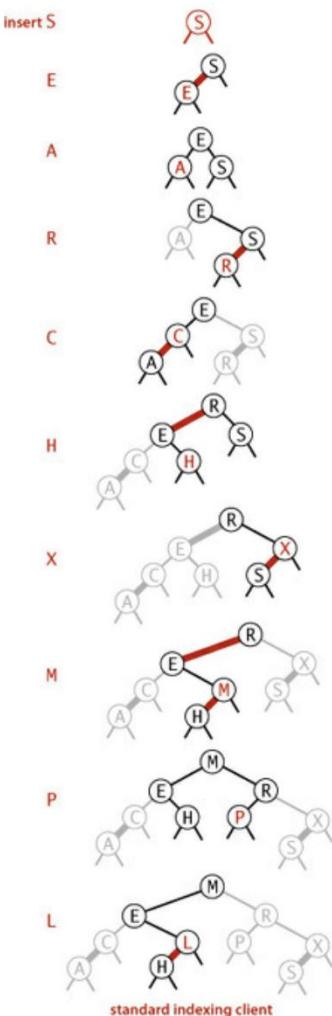
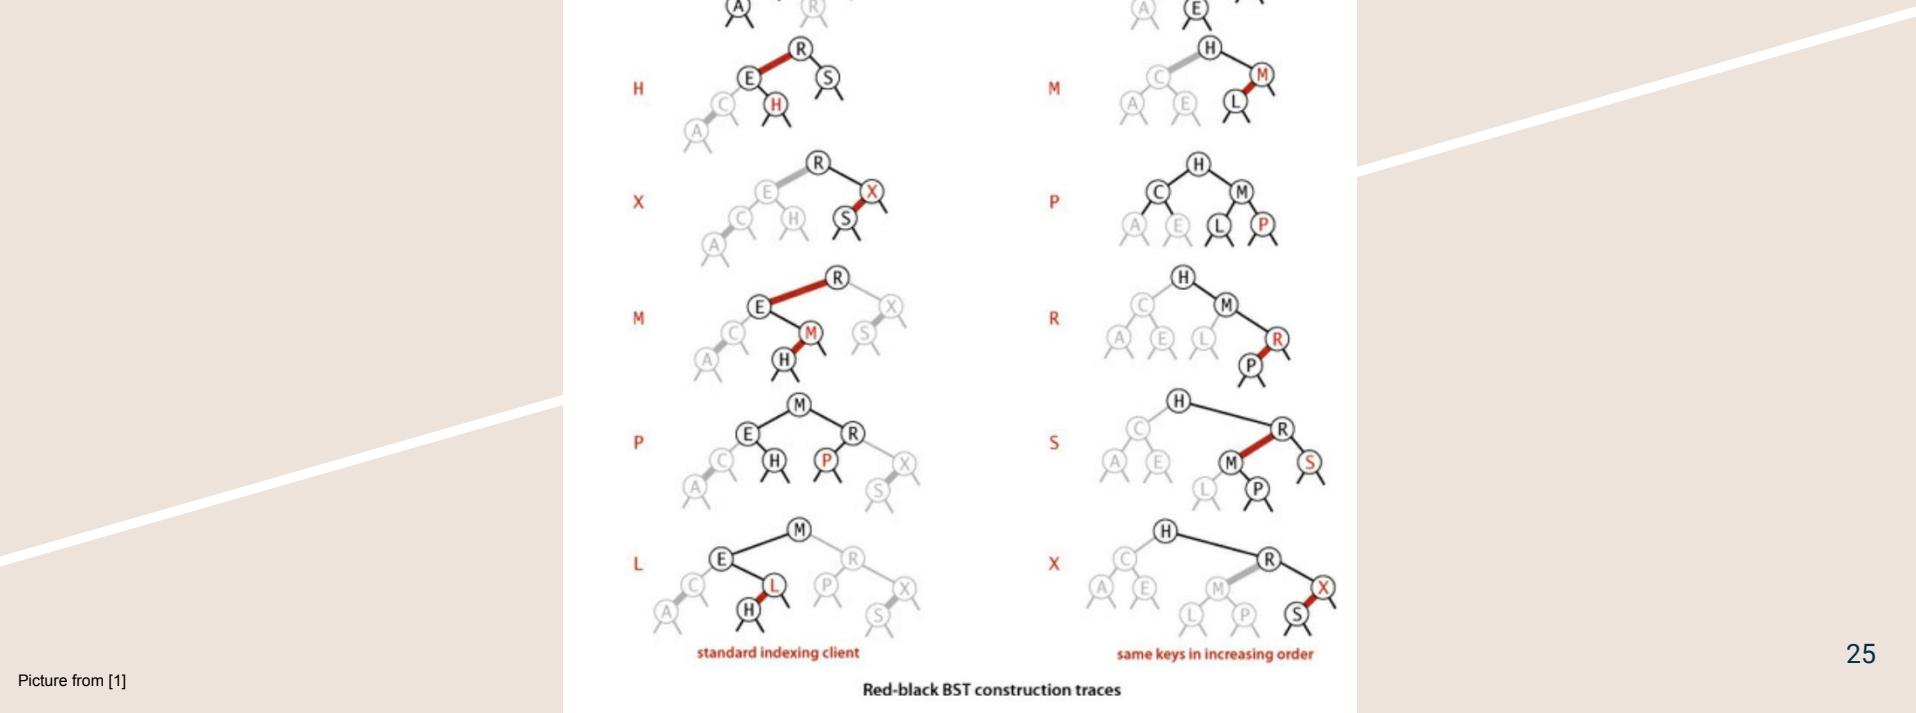
Case 4:  $X.left$  is red,  $X.left.left$  is red.



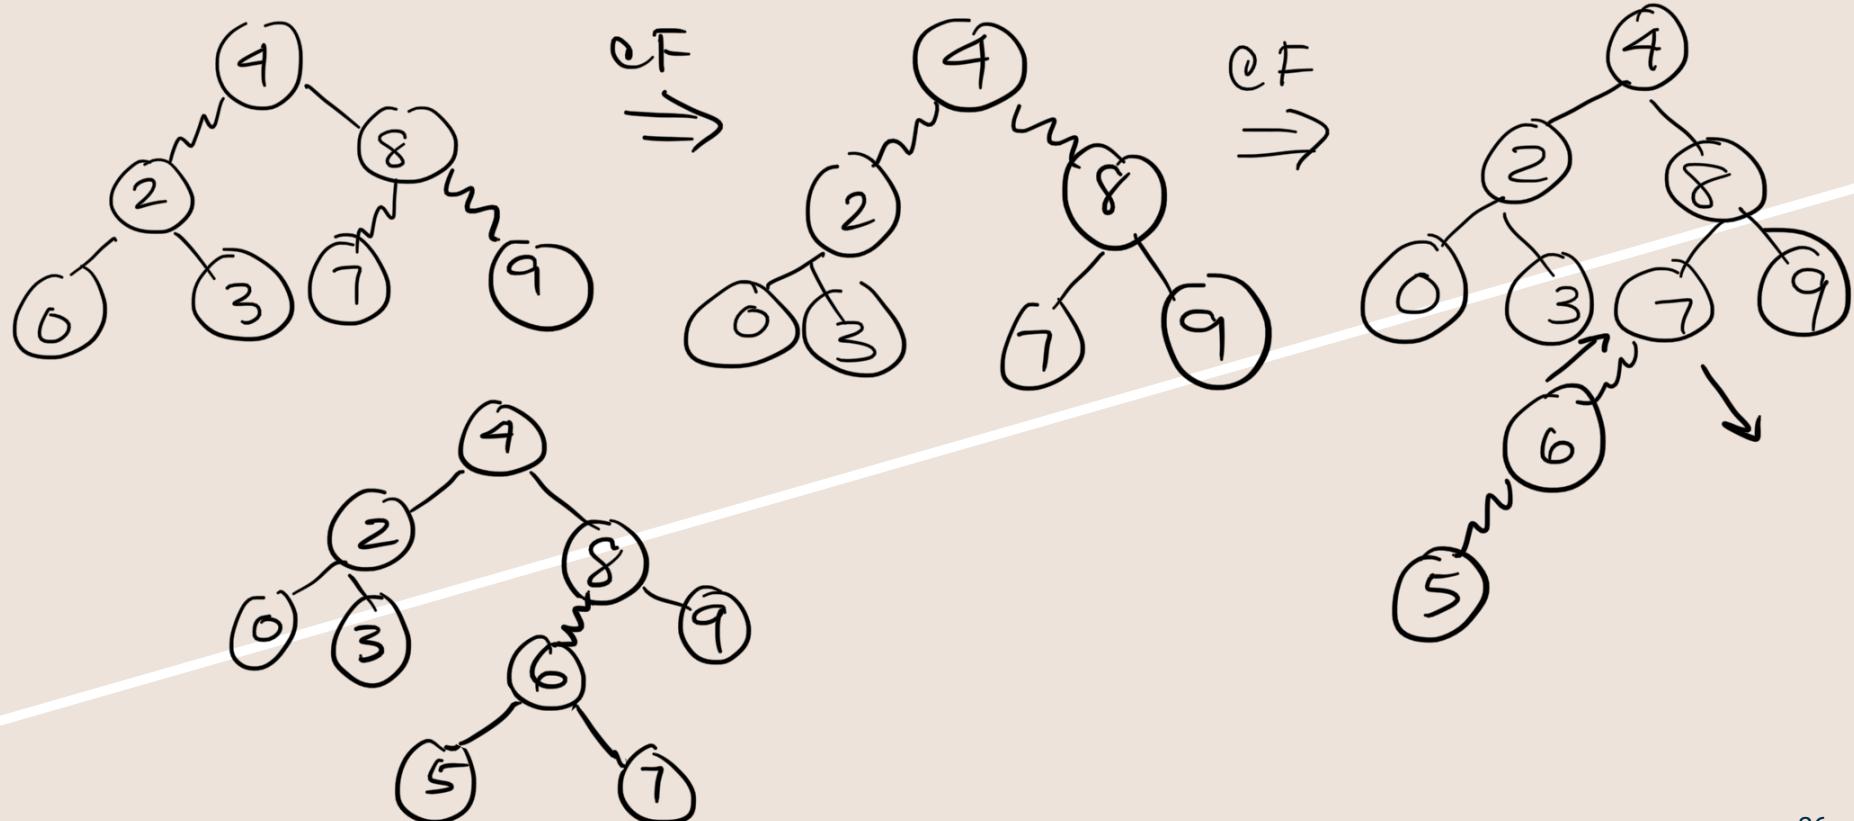
This case is only possible with *delete*.

Case 5:  $X.right$  is red,  $X.right.left$  is red.





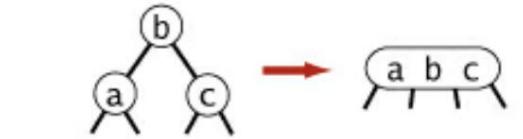
~~Example: insert 0, insert 4, insert 7, insert 2, insert 3,  
insert 8, insert 9, insert 6, insert 5~~



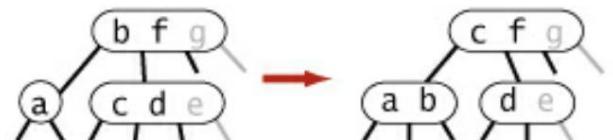
# delete

- It's easier to delete a red node than a black node. Why?
- Note: The computer won't know if a node to be deleted is red or black until it finds that node.
- Key Idea: Use the inverses of the transformations discussed previously to "carry a red link down the line"--just in case.
- How is this related to deleting from a 2-3 tree?

at the root



on the way down



at the bottom



Transformations for delete the minimum

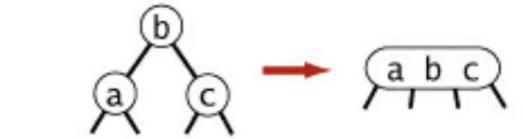
## Big Idea for Deletion:

It is easy to delete a key from a 3-node (or a 4-node) at the bottom of the tree. The tricky part is deleting a 2-node. (WHY?)

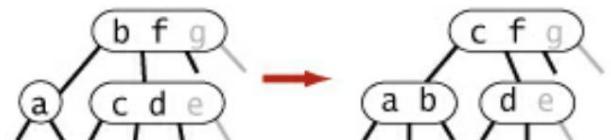
So we can transform the tree on the way down to ensure that the current node is not a 2-node. (HOW?)

What is the equivalent idea for a red-black tree?

at the root



on the way down



at the bottom



Transformations for delete the minimum

## Big Idea for Deletion:

It is easy to delete a key from a 3-node (or a 4-node) at the bottom of the tree. The tricky part is deleting a 2-node. (WHY?)

So we can transform the tree on the way down to ensure that the current node is not a 2-node. (HOW?)

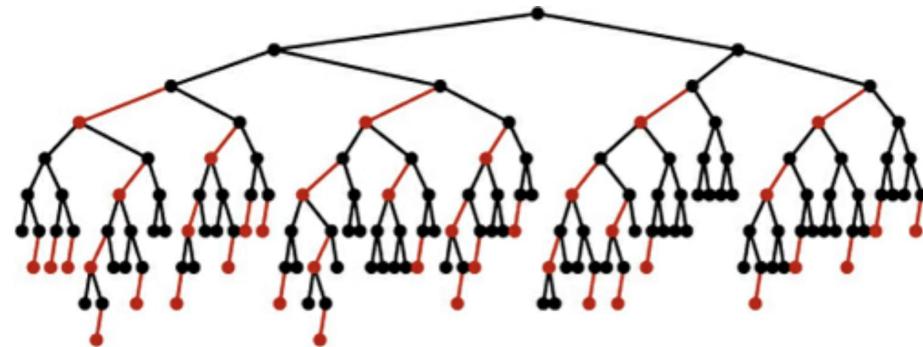
What is the equivalent idea for a red-black tree?

Make sure the current node is always red.

~~Example: delete 1, delete 6, delete 2, delete 7, delete 4,~~  
~~delete 3, delete 5~~

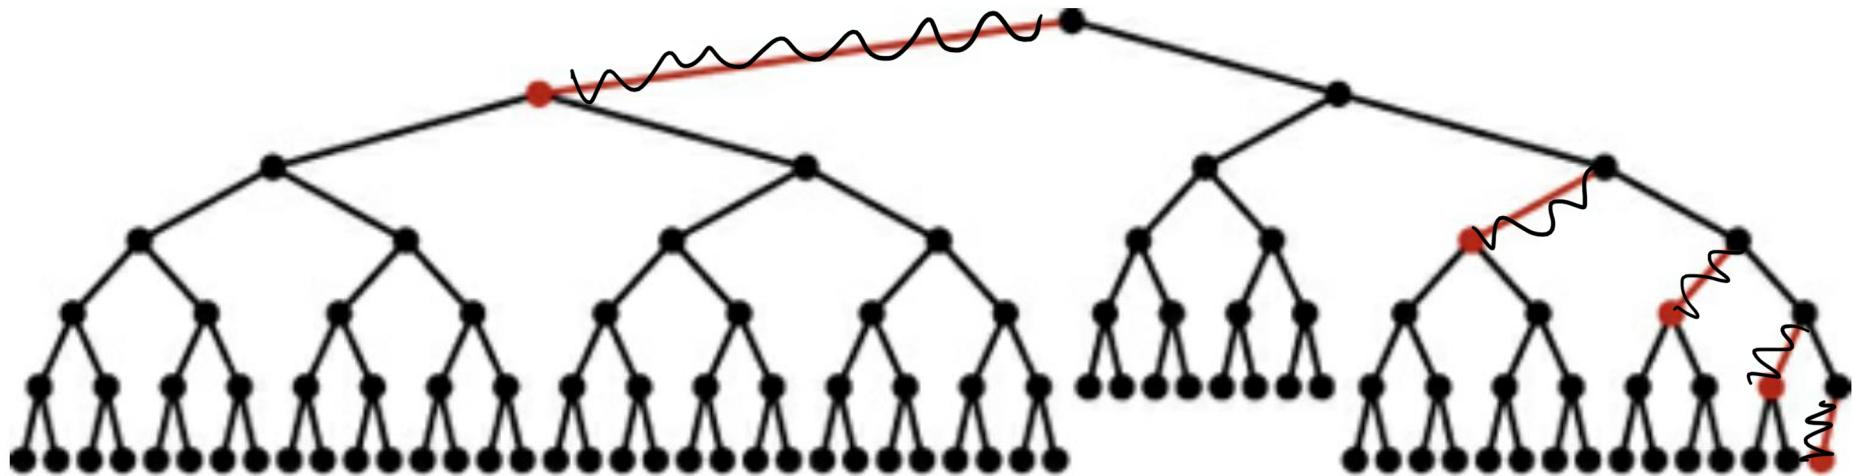
# Properties of Red Black Trees

- **Near** perfect balance
- 1-1 correspondence with 2-3 Trees
- Height is no more than  $2\lg N$   
(worst case is all 2-nodes except the leftmost path, which is 3 nodes and so twice as long)
- Average length of root to a node is  $\sim 1.00\lg N$



Typical red-black BST built from random keys (null links omitted)

Picture from [1]



Red-black BST built from ascending keys (null links omitted)

---

**Proposition I.** In a red-black BST, the following operations take logarithmic time in the worst case: search, insertion, finding the minimum, finding the maximum, floor, ceiling, rank, select, delete the minimum, delete the maximum, delete, and range count.

algorithm (data structure)	worst-case cost (after N inserts)		average-case cost (after N random inserts)		efficiently support ordered operations?
	search	insert	search hit	insert	
<i>sequential search (unordered linked list)</i>	N	N	N/2	N	no
<i>binary search (ordered array)</i>	$\lg N$	N	$\lg N$	N/2	yes
<i>binary tree search (BST)</i>	N	N	$1.39 \lg N$	$1.39 \lg N$	yes
<i>2-3 tree search (red-black BST)</i>	$2 \lg N$	$2 \lg N$	$1.00 \lg N$	$1.00 \lg N$	yes

**Cost summary for symbol-table implementations (updated)**

# References

- [1] *Algorithms, Fourth Edition*; Robert Sedgewick and Kevin Wayne (and associated slides)
- [2] Slides from <https://www.cs.princeton.edu/~rs/talks/LLRB/RedBlack.pdf>