

Skiplists

A probabilistic data structure...

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

The Skiplist...

- ...does not *guarantee* $O(\log N)$ runtime for operations,...
- ...does provide $O(\log N)$ performance *with a high probability*.
- ...can provide a balance between good performance and ease of implementation.

Ideally...

- You would have enough levels so that the ideal search would skip half the values, then a quarter of the remaining values, then an eighth of the remaining values, etc...
- Which would mean that a *search* (or an *insert*) would be $O(\log N)$ with high probability.
- Space: If the head node isn't counted,
 - $\frac{1}{2}$ the nodes should have 1 pointer...
 - $\frac{1}{4}$ should have 2 pointers...
 - $\frac{1}{8}$ should have 3 pointers...
 - and so on...
- However, we don't enforce the *ideal* situation.
- Instead, we rely on probability and randomization.

Inserting & removing nodes...

- Instead, a new node is assigned a level at random where there is
 - a 50% chance that it will have 1 pointer (the lowest level),
 - a 25% chance that it will have 2 pointers (the second lowest level)...
 - and so on...
- This is what makes skiplists *probabilistic*.
- <https://opendsa-server.cs.vt.edu/ODSA/Books/CS3/html/SkipList.html#id1>


$$N + \frac{N}{2} + \frac{N}{4} + \dots + 1 \Rightarrow O(N)$$

Analysis

- Consider the extremes:



- All the nodes are at a high level, so lots of pointers. In that case, the insert cost could be quite high (even $O(N)$). (Basically, this would be a linked list with extra pointers).
 $(\frac{1}{2}) (\frac{1}{2}) (\frac{1}{2}) \dots (\frac{1}{2})$
- All the nodes are at a low level (e.g. level 0), so they only have 1 pointer each, which is basically just a linked list. In that case, insert and search could be $O(N)$.

$\frac{1}{2^n}$

- The good news: These extremes are unlikely to happen. For example, there is only $\frac{1}{1024}$ chance that 10 nodes in a row will be in level 0.
- Similar to Quicksort, we accept that the randomization will make sure that the worst case is unlikely to happen.
- Expected memory use: $O(N)$

Skiplist vs. BST

- Recall that a BST made from random keys is expected to have a height of approximately $\log N$.
- A skiplist's performance is expected to be about the same as a BST in which the keys are inserted in random order.
- The key difference: There is no randomization built into a BST.
- In fact: The larger the skiplist gets, the closer it will get to $O(\log N)$ performance because it becomes less and less likely that the worst case will happen.